

Assingment 1 : Deep Learning for CV - Building a CNN

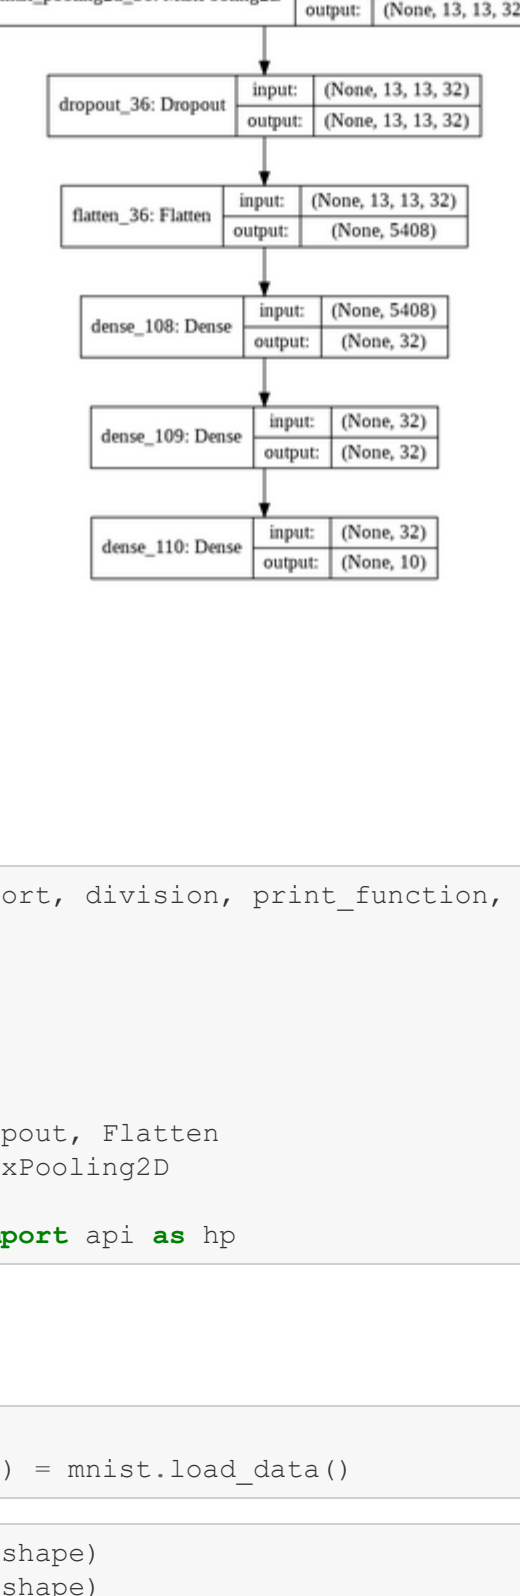
Index Number: 17000475

Dataset- MNIST

Aim- Classify digit images from 0 to 9

1. Architecture of the CNN Model

```
In [ ]: from keras.utils.vis_utils import plot_model
plot_model(model, to_file='model_plot.png', show_shapes=True, show_layer_names=True)
```



2. Source Code

Import Modules

```
In [ ]: from future import absolute_import, division, print_function, unicode_literals
import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
import datetime
from tensorflow import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.utils import np_utils
from tensorflow.keras.callbacks import TensorBoard
```


Dataset Preparation and Inspect

```
In [ ]: mnist = keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [ ]: print("train image shape:", x_train.shape)
print("train label shape:", y_train.shape)
print("test image shape:", x_test.shape)
print("test label shape:", y_test.shape)
```

```
train image shape: (60000, 28, 28)
train label shape: (60000,)
test image shape: (10000, 28, 28)
test label shape: (10000,)
```

```
In [ ]: plt.figure()
plt.imshow(x_train[1])
# plt.colorbar()
# plt.grid(False)
plt.show()
```



```
In [ ]: fig, axes = plt.subplots(3, 6, figsize=(8,8))
fig.subplots_adjust(hspace=0.5, wspace=0.5)
for i, ax in enumerate(axes.flat):
    ax.imshow(x_train[i])
    ax.set_xticks([])
    ax.set_yticks([])
plt.show()
```

```
In [ ]: #Normalization - Rescaling pixels on scale 0 to 1
x_train = x_train / 255.0
x_test = x_test / 255.0
```

Tensorboard Setup

```
In [ ]: #initializing Tensorboard extension
load_ext tensorboard
#creating a logs directory
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
In [ ]: # Clear any logs from previous runs
!rm -rf ./logs/
```

```
In [ ]: x_train.shape[1:]
```

```
Out [ ]: (28, 28, 1)
```

Hyperparameter Selection

```
In [ ]: HP_NUM_UNITS_1 = hp.HParam('num_units_1', hp.Discrete([16,32,64]))
HP_NUM_UNITS_2 = hp.HParam('num_units_2', hp.Discrete([16,32,64]))
HP_NUM_UNITS_3 = hp.HParam('num_units_3', hp.Discrete([16,32,64]))
HP_DROPOUT = hp.HParam('dropout', hp.RealInterval(0.1, 0.2))
HP_OPTIMIZER = hp.HParam('optimizer', hp.Discrete(['adam']))
```

METRIC_ACCURACY = 'accuracy'

```
with tf.summary.create_file_writer('logs/hparam_tuning') as default():
    hp.hparams.config(
        hparams=[HP_NUM_UNITS_1, HP_NUM_UNITS_2, HP_DROPOUT, HP_NUM_UNITS_3, HP_OPTIMIZER],
        metrics=[HP_Metric(METRIC_ACCURACY, display_name='Accuracy')],
    )
```

```
In [ ]: def train_test_model(hparams):
```

```
    # number of possible label values
    model = Sequential()
    #input_shape = (height,width,channels)
    # Here Batch size will be None at the begining at later when we fit the model it will be used
    model.add(Conv2D(hparams[HP_NUM_UNITS_1], kernel_size=(3,3), activation='relu', padding='same', input_shape=(28, 28, 1)))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(hparams[HP_DROPOUT]))
    #flattened 3d output to 1d
    model.add(Flatten())
    model.add(Dense(hparams[HP_NUM_UNITS_2], activation='relu'))
    model.add(Dense(hparams[HP_NUM_UNITS_3], activation='relu'))
    # Output classes =10 Hence used 10 outputs in Dense layer
    model.add(Dense(10, activation='softmax'))

    model.compile(
        optimizer=hparams[HP_OPTIMIZER],
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'],
    )
    model.fit(x_train, y_train, epochs=3)
    _, accuracy = model.evaluate(x_test, y_test)
    return accuracy
```

```
In [ ]: def run(run_dir, hparams):
```

```
    with tf.summary.create_file_writer(run_dir) as default():
        hp.hparams.config(
            hparams=[HP_NUM_UNITS_1, HP_NUM_UNITS_2, HP_DROPOUT, HP_NUM_UNITS_3, HP_OPTIMIZER],
            metrics=[HP_Metric(METRIC_ACCURACY, accuracy, step=1)]
        )
```

```
In [1]: session_num = 0
```


```
for num_units_1 in HP_NUM_UNITS_1.domain.values:
    for num_units_2 in HP_NUM_UNITS_2.domain.values:
        for dropout_rate in HP_DROPOUT.domain.max_value, HP_DROPOUT.domain.min_value:
            for num_units_3 in HP_NUM_UNITS_3.domain.values:
                hparams = {
                    HP_NUM_UNITS_1: num_units_1,
                    HP_NUM_UNITS_2: num_units_2,
                    HP_DROPOUT: dropout_rate,
                    HP_NUM_UNITS_3: num_units_3,
                    HP_OPTIMIZER: 'adam',
                }
                run_name = "run-{} {}".format(session_num, hparams)
                print(h_name, hparams[h] for h in hparams)
                run('logs/hparam_tuning/' + run_name, hparams)
                session_num += 1
```

Part of training process

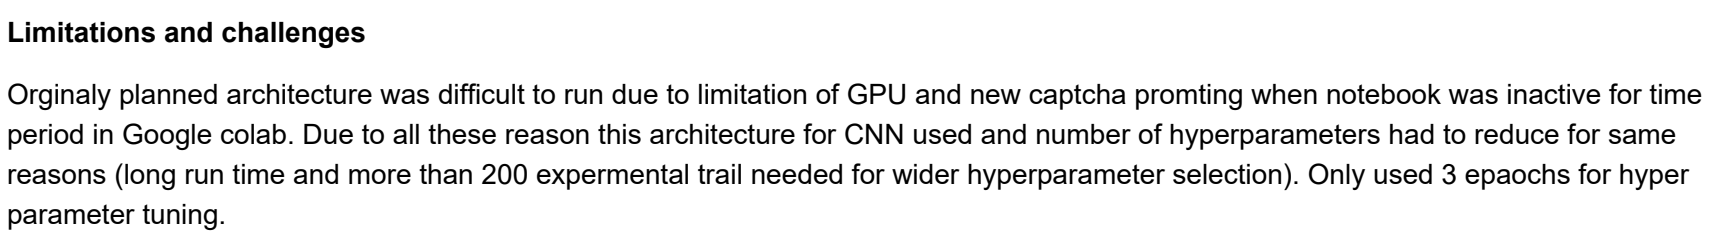
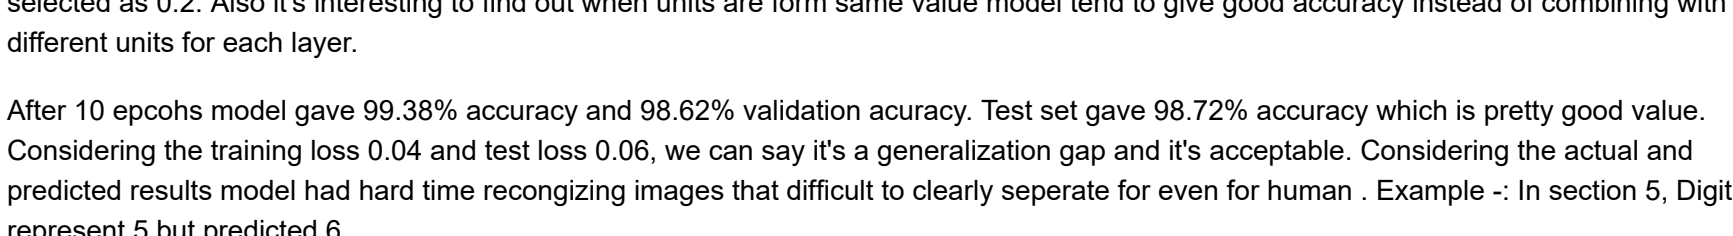




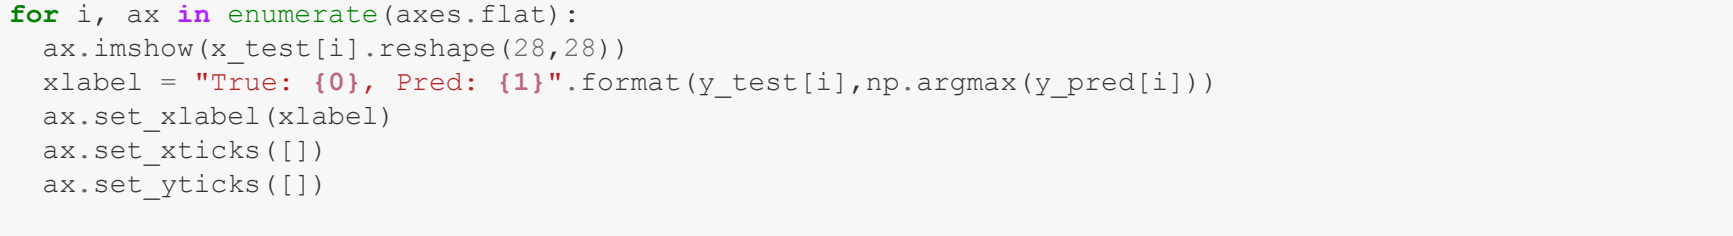
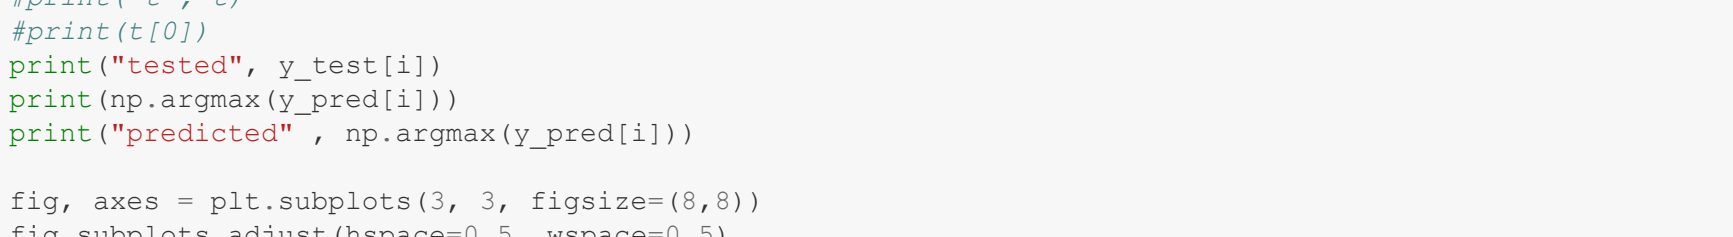











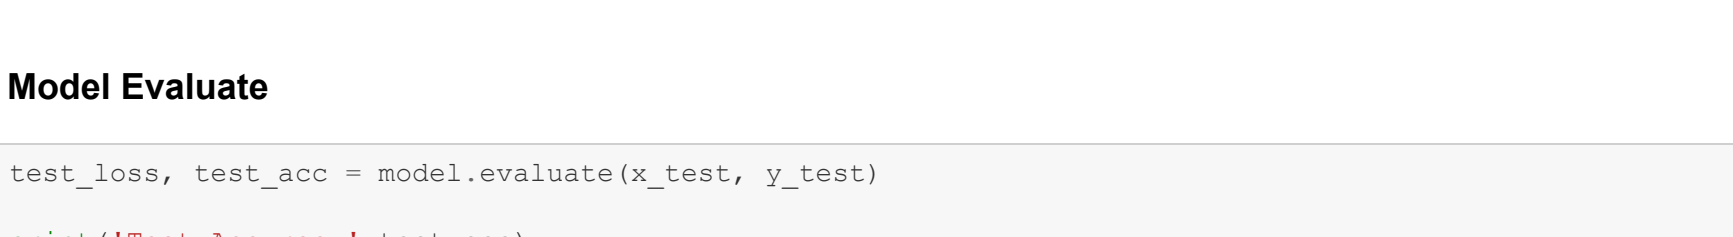









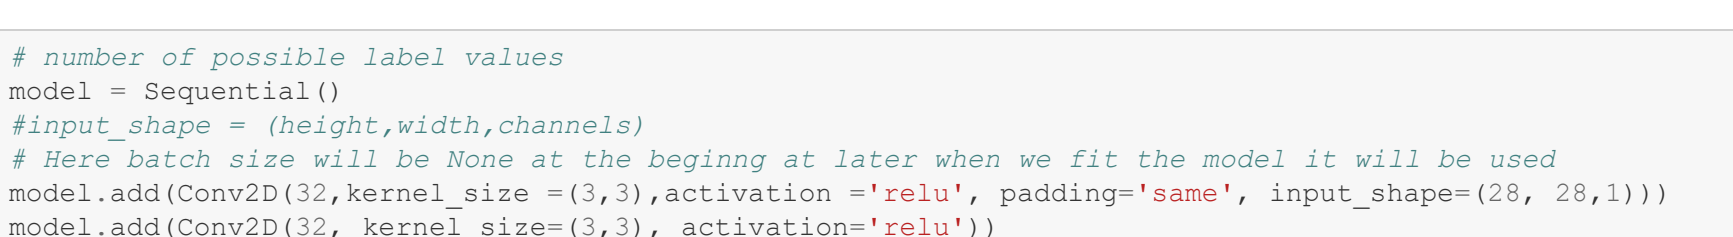
```
--- Starting trial: run-34
('num_units_1': 64, 'num_units_2': 64, 'dropout': 0.2, 'num_units_3': 32, 'optimizer': 'adam')
Epoch 1/3
1875/1875 [=====] - 166s 89ms/step - loss: 0.3052 - accuracy: 0.9016
Epoch 2/3
1875/1875 [=====] - 165s 89ms/step - loss: 0.0511 - accuracy: 0.9938
Epoch 3/3
1875/1875 [=====] - 165s 89ms/step - loss: 0.0339 - accuracy: 0.9991
--- Starting trial: run-35
('num_units_1': 64, 'num_units_2': 64, 'dropout': 0.2, 'num_units_3': 64, 'optimizer': 'adam')
Epoch 1/3
1875/1875 [=====] - 169s 89ms/step - loss: 0.2786 - accuracy: 0.9126
Epoch 2/3
1875/1875 [=====] - 168s 89ms/step - loss: 0.0462 - accuracy: 0.9853
Epoch 3/3
1875/1875 [=====] - 167s 89ms/step - loss: 0.0306 - accuracy: 0.9898
1875/1875 [=====] - 7s 22ms/step - loss: 0.0376 - accuracy: 0.9888
```

```
In [2]: !tensorboard --logdir logs/hparam_tuning
```

3. Justification for Hyperparameter selection



Trial ID	num_units_1	num_units_2	dropout	num_units_3	optimizer	Accuracy
0f006d55ca24...	32.0000	32.0000	0.200000	32.0000	adam	0.98910




Other Images

conv2d_72

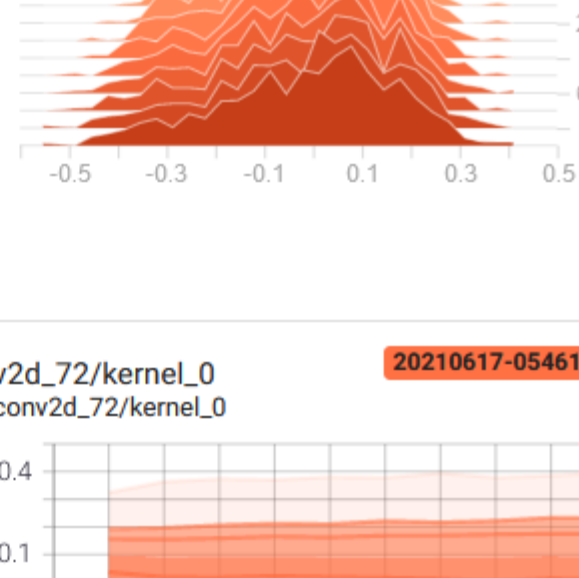
conv2d_72/bias_0
tag: conv2d_72/bias_0

20210617-054616/train



conv2d_72/kernel_0
tag: conv2d_72/kernel_0


20210617-054616/train



conv2d_72


conv2d_72/bias_0
tag: conv2d_72/bias_0

20210617-054616/train



conv2d_72/kernel_0
tag: conv2d_72/kernel_0

20210617-054616/train



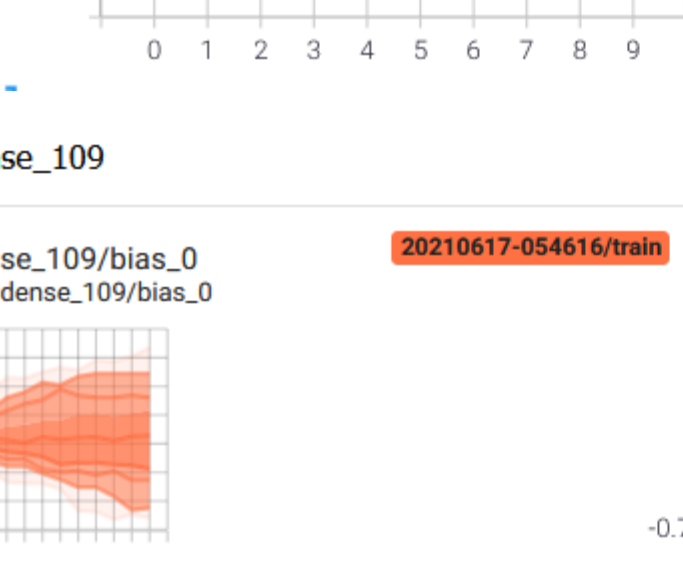
--

--

conv2d_73

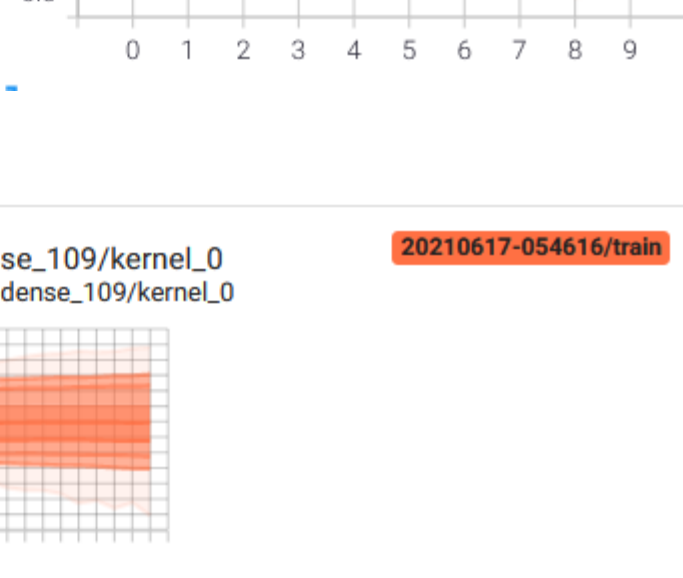
conv2d_73/bias_0
tag: conv2d_73/bias_0

20210617-054616/train



conv2d_73/kernel_0
tag: conv2d_73/kernel_0

20210617-054616/train




--

--

dense_108


dense_108/bias_0
tag: dense_108/bias_0

20210617-054616/train



dense_108/kernel_0
tag: dense_108/kernel_0

20210617-054616/train



--

--

dense_109

dense_109/bias_0
tag: dense_109/bias_0

20210617-054616/train



dense_109/kernel_0
tag: dense_109/kernel_0

20210617-054616/train



--

--

dense_110

dense_110/bias_0
tag: dense_110/bias_0

20210617-054616/train



dense_110/kernel_0
tag: dense_110/kernel_0

20210617-054616/train

